

Fall 2021
ECE 5770: GPU Accelerated Computing
Prof. Shadi G. Alawneh
Deadline: Tue Nov 2 11:59 pm

Assignment 3

Histogram

Description

The purpose of this assignment is to implement an efficient histogramming algorithm (using privatization technique) for an input array of integers within a given range. Each integer will map into a single bin, so the values will range from 0 to (NUM_BINS - 1). The histogram bins will use unsigned 32-bit counters that must be saturated at 127 (i.e. no roll back to 0 allowed). The input length can be assumed to be less than 2^{32} . NUM_BINS is fixed at 4096 for this assignment. This can be split into two kernels: one that does a histogram without saturation, and a final kernel that cleans up the bins if they are too large. These two stages can also be combined into a single kernel.

Instructions

Edit the code to perform the following:

- Allocate device memory
- Copy host memory to device
- Initialize thread block and kernel grid dimensions
- Invoke CUDA kernel
- Copy results from device to host
- Free device memory
- Write the CUDA kernel

Instructions about where to place each part of the code is demarcated by the `//@@@` comment lines.

Local Setup Instructions

Requirements

libwb:

<https://github.com/abduld/libwb>

Compile instructions (Unix)

You will need libwb (a library by the authors of the teaching kit) for all assignments in this course.

The easiest way to get everything working would be to create a directory for class assignments, and place libwb in that directory alongside the individual assignment directories.

To build libwb:

```
cd [path_to_libwb]
```

```
mkdir build
```

```
cd build
```

```
cmake ..
```

```
make
```

To build Histogram:

```
cd [path_to_Histogram]
```

```
mkdir build
```

```
cd build
```

```
cmake ..
```

```
make
```

This should just work if the Histogram and libwb directories are in the same parent directory.

Testing

The "Dataset" directory contains some sample input / output data to test your implementation. To test a particular input, run (from the Histogram root directory):

```
./build/Histogram -i Dataset/##/input.raw -e Dataset/##/output.raw  
-o output.raw -t integral_vector
```

This will produce some output, part of which should indicate whether or not your output matches the expected output.

To test all 6 examples at once, run (again from the Histogram root directory):

```
python test.py
```

Upon completion, this should tell you how many of the samples you produced the correct output for, and indicate which ones (if any) were incorrect.

Submission

You can simply submit your solution on Moodle. There is no need to include the dataset or any build files, just the Histogram.cu is fine!